

Implementation and Analysis of BST, AVL and RBT Data Structures

Sebastiano Babich (172249) - Leonardo Mazzon (172587)
Damiano Netto (171948)

University of Udine, Academic Year 2025 - 2026

Contents

1	Abstract	3
2	Introduction	3
3	Theoretical Background	3
3.1	Binary Search Tree (BST)	4
3.2	Adelson-Velsky Landis Tree (AVL)	4
3.3	Red Black Tree (RBT)	5
4	Implementation	6
5	Test Cases Setup	8
5.1	Choosing appropriate input sizes	8
5.2	Choosing number of test cases	8
6	Results and Discussion	10
6.1	Insertion Tests	10

6.2	Rotations Count Tests	10
6.3	Heights Tests	10
7	Conclusions	11
7.1	Final Remarks	11

1 Abstract

The goal of the project is to implement and compare some of the data structures studied in the “Algorithms and Data Structures” course at the University of Udine, specifically standard and self-balancing binary search trees. The project focuses on analyzing the time complexity of insertion and deletion within these structures. We also analyze the heights and rotation performed by the trees of our implementation.

2 Introduction

Binary search trees (BST) are among the most widely used data structures, since they yield logarithmic asymptotic time in the average case of search, insertion and deletion. As the name suggests, each node in the tree has to have two children, which can either be another node or a null pointer, and the tree’s height is the number of edges in the longest path between the root of the tree and a leaf. Two examples of Red-Black Trees (balanced variation of BSTs, covered in our analysis) being used are `std::set`[4] in C++’s standard library and the `Completely Fair Scheduler`[2] in Linux.

3 Theoretical Background

In all tree data structures, height is a fundamental property: keeping its value as small as possible helps ensure efficient operations thanks to more efficient search, insertion and deletion operations.

Trees can be divided into two categories based on how they manage their height: **binary search trees** and **balanced binary trees**. Standard binary search trees do not actively control their height, while balanced binary trees use specific rules and rotation operations to keep the height as small as possible.

Trees of the latter kind are designed to maintain a height close to the order of $\log_2(n)$, where n is the number of nodes inserted into the tree.

In this project, we studied three types of trees, two of which implement self-balancing techniques.

3.1 Binary Search Tree (BST)

Binary search trees do not impose any restrictions on their height. Their purpose is to provide efficient search, insertion and deletion operations maintaining an ordered structure of the stored elements. However, since they do not include any mechanism to balance them, the tree can become unbalanced depending on the order in which nodes are inserted. If for example, nodes are inserted in an orderly manner, it will result in the Worst-Case complexity (which is listed below in asymptotic time), making the tree isomorphic to a linked list. So inserting values in a “random” order will yield better results.

Operation	Average Case	Worst Case
Search	$\mathcal{O}(\log n)$	$\mathcal{O}(n)$
Insert	$\mathcal{O}(\log n)$	$\mathcal{O}(n)$
Delete	$\mathcal{O}(\log n)$	$\mathcal{O}(n)$

Table 1: Time complexity in BST operations

3.2 Adelson-Velsky Landis Tree (AVL)

AVL trees are a data structure that implements self-balancing binary search trees. The search operation in the tree exactly mirrors the search in BST. Insertion initially uses the BST insertion method, with the difference that a self-balancing rotation is performed afterwards.

Specifically, this operation aims to identify imbalances in the tree by comparing the heights of two child nodes. Specifically, the goal is to keep total height as low as possible, having an upper bound of $1.44 \cdot \log_2(n + 2)$ (Adelson-Velsky and Landis [1]). This yields better time complexity in search operations, however, its main drawback is insertion time, as rotations provide a lot of overhead, but this is necessary to achieve the aforementioned height.

Operation	Average Case	Worst Case
Search	$\mathcal{O}(\log n)$	$\mathcal{O}(\log n)$
Insert	$\mathcal{O}(\log n)$	$\mathcal{O}(\log n)$
Delete	$\mathcal{O}(\log n)$	$\mathcal{O}(\log n)$

Table 2: Time complexity in AVL operations

3.3 Red Black Tree (RBT)

An RBT is a self-balancing BST where every node is colored as red or black and satisfies certain conditions. So for each node we also have to keep track to its color. If T is a RBT, for all n node of T , we call $color(n)$ its color.

Let T a BST where each node has a color value associated ($color(n)$).

T is a RBT $\iff T$ satisfies the following conditions:

- For all n node in T , $color(n)$ must be either BLACK or RED (it must be a total function from the set of all nodes of T to $\{\text{BLACK}, \text{RED}\}$)
- NIL leaves are considered black.
- For all red nodes n in T , n must not have a red child.
- Every path from a n node in T to its descendant NIL leaves must contain the same number of black nodes.

For readonly operations (like search, inorder visit, etc...) we have the same complexity as BST. Instead, insertion and removal operations are augmented using two new helper functions whose main purpose is to maintain the properties of the RBT, using rotations like AVL Trees.

These trees are a compromise between the aforementioned trees, as they yield efficient results even in insertion and deletion, due to their limited use of rotations during the balancing of the tree. Although this means that the tree might not achieve is absolute minimum height, as it has an upper bound of $2 \cdot \log_2(n + 1)$ (Cormen et al. [3, p. 332]), which is slightly worse than the height of the AVL. Since this tree combines fast insertion akin to a BST and a search almost as optimal as an AVL. It is suitable in systems requiring frequent search and insert/delete operation

Operation	Average Case	Worst Case
Search	$\mathcal{O}(\log n)$	$\mathcal{O}(\log n)$
Insert	$\mathcal{O}(\log n)$	$\mathcal{O}(\log n)$
Delete	$\mathcal{O}(\log n)$	$\mathcal{O}(\log n)$

Table 3: Time complexity in RBT operations

4 Implementation

The data structures (BST, AVL and RBT) were implemented as separate modules in the project, with similar interfaces for insertion, deletion, and search. Each tree maintains additional information useful for analysis, such as current height and number of rotations. Balancing operations were separated into dedicated functions within each type of tree to comply with the rules imposed by the data structure used.

These are the main functionalities implemented for each tree:

- **BST (Binary Search Trees)**

- **Node implementation:** a *Node* class has been implemented to store the key, the references to the left and right children, and a reference to the parent node. Each one of these could be another node, or null reference.
- **Node insertion and deletion:** insertion and deletion operations are implemented according the standard algorithms of Binary Search Trees, maintaining the BST ordering property.
- **Search operations:** efficient search, minimum, maximum, predecessor and successor operations are provided by exploiting the ordering property of the Binary Search Tree.
- **Tree traversals operations:** inorder, preorder and postorder are implemented iteratively to permit the visit of each tree created.
- **Tree rotations:** left and right rotation operations are implemented as fundamental tree transformations. These operations are reused by the balanced tree implementations and a rotation counter is kept to be used in the experimental analysis.

- **AVL (Adelson-Velsky Landis Trees)**

- **AVLNode implementation:** an *AVLNode* class, extending the *Node* class contained in BST, has been created. Each node stores an additional *height* field, which is required to compute the balance factor.
- **Balance factor computation:** the balance factor of each node is calculated as the difference between the heights of its left and right subtrees. That value is used to detect unbalanced configurations.
- **Height management:** the height of each node is updated after insertions, deletions and rotations through the *update_height* function, avoiding the need of recomputing heights recursively.

- **Node insertion and deletion:** insertion and deletion operations are implemented by extending the BST algorithms and performing the required rebalancing after each update, checking the `balance_factor` calculated. The operations which can occur can be rotations in left or right direction.
 - **Tree rotations:** rotations are used as balancing operations to restore the AVL property after insertions and deletions. These operations are partially inherited from the BST class. Single and double rotations are selected according to the imbalance configuration detected by the balance factor.
- **RBT (Red-Black Trees)**
 - **RBTNode implementation:** *RBTNode* class has been created to extend *Node* class contained in BST. Each node in fact stores an additional *color* field, which can be either red or black, fundamental to maintain the red-black tree properties.
 - **Color management:** utility functions have been implemented to get, set and invert node colors. Missing children are considered black, simplifying the handling of boundary cases during insertion and deletion.
 - **Node insertion and deletion:** insertion and deletion are implemented by extending the standard BST functions. After each insertion or deletion, a rebalancing procedure is performed, involving recoloring operations and tree rotations. The required operations restore the red-black properties while preserving the BST ordering.
 - **Tree rotations:** left and right rotations inherited from BST are reused during the balancing procedures to modify the tree structure while preserving the ordering property.

5 Test Cases Setup

5.1 Choosing appropriate input sizes

The first issue to manage during the implementation of the test cases was the choice of the input sizes on which the experiments would be performed. We chose a wide range of tree sizes to make the simulation as realistic as possible, highlighting the performance differences among the different data structures.

A valid range of values for the number of nodes had to allow the analysis in multiple orders of magnitude. Instead of using a linear distribution of input sizes, which could have caused issues in tests on large number of elements, a geometric progression was adopted in order to obtain a more uniform distribution on a logarithmic scale.

This approach allowed us to perform a better analysis of the growth trends of the different tree implementations.

These were the parameters selected:

$$n_{min} = 1000 \qquad n_{max} = 1\,000\,000 \qquad samples = 100$$

The list of input sizes is generated through the *calc_lista_val_n* function, which computes the values of n using a geometric progression. The common ratio of the progression is defined as:

$$n_i = \lfloor n_{min} \cdot c^i \rfloor \qquad c = \left(\frac{n_{max}}{n_{min}} \right)^{\left(\frac{1}{samples-1} \right)}$$

The geometric progression produces a total of 100 samples, gradually increasing from the minimum value of nodes up to the maximum value. The first and last values of the sequence are reported below:

$$1000, 1072, 1150, 1233, 1322, \dots, 811\,131, 869\,749, 932\,603, 1\,000\,000$$

Then for each n_i , we need to choose how many test cases we want to make.

5.2 Choosing number of test cases

At the start we tried a constant number of test cases for each n_i . But then we quickly realized that this didn't make much sense because if the constant is too small, then

for larger n_i we wouldn't cover a good amount of cases, but if the constant is too big, then for smaller n_i we would make a lot of redundant test cases wasting time and resources.

So we changed it to be this function:

$$f(n) = \min + k\sqrt{n} \tag{1}$$

The reason why we choose $\sqrt{n}$ is that it does not grow so big like n that it will make the testing very inefficient (think when n is 1000000) but it is still big enough to cover a lot of cases (on $n = 1000000$, 1000 cases at least if $k \geq 1$). The $\min$ is needed so that we do at least $\min$ tests for small n , and we choose k if we want to do more or less tests.

In the tests we choose $\min = 100$ and $k = 2$ as it was a good compromise.

So the number of tests for each n_i is $100 + 2\sqrt{n_i}$.

6 Results and Discussion

The obtained results correctly reflected the theoretical concepts studied.

6.1 Insertion Tests




Figure 1: AVL insertion test

6.2 Rotations Count Tests

6.3 Heights Tests

7 Conclusions

To conclude this project, we compared all the data structures we analyzed and drew conclusions regarding the ideal use cases for each type of tree.

7.1 Final Remarks

The comparison demonstrates that the additional complexity required to implement self-balancing trees is justified by the significant longterm gains in efficiency and robustness offered by these data structures. While a standard BST provides a simpler implementation and execution, it remains vulnerable to worst-case scenarios making it less reliable for large-scale datasets. Instead, choosing between an AVL and a RBT allows us to optimize the system based on whether the specific application requires fast searches or a balanced mix of frequent updates and lookups.

Bibliography & Sitography

- [1] Georgy Adelson-Velsky and Evgenii Landis. “An Algorithm for the Organization of Information”. In: (1962).
- [2] *CFS Scheduler - The Linux Kernel documentation*. Archived on [archive.org](https://web.archive.org/web/20260424153328/https://www.kernel.org/doc/html/latest/scheduler/sched-design-CFS.html#the-rbtree) on April 24th, 2026. URL: <https://web.archive.org/web/20260424153328/https://www.kernel.org/doc/html/latest/scheduler/sched-design-CFS.html#the-rbtree>.
- [3] Thomas H. Cormen et al. *Introduction to Algorithms*. 4th ed. MIT Press, 2022.
- [4] *std::set - cppreference.com*. Archived on [archive.org](https://web.archive.org/web/20260708070457/https://en.cppreference.com/cpp/container/set) on July 8th, 2026. URL: <https://web.archive.org/web/20260708070457/https://en.cppreference.com/cpp/container/set>.